

Master of Science in Engineering with Specialisation in Data Science

**Master Thesis**

# **The Effect of Dynamically Built Lateral Connection in Deep Learning Systems**

**On Bridging the Gap between Neuroscience and Deep Learning**

Pascal Sager

August 20, 2022

Zurich University of Applied Sciences



Dedicated to the dedicated.

– Pascal Sager



*"Max Planck said, 'Science progresses one funeral at a time.'  
The future depends on some graduate student who is deeply  
suspicious of everything I have said."*

– Geoffrey Hinton, University of Toronto, 2017.



# **Zusammenfassung**





# Abstract



# Preface

This thesis is the last step stone before I will hold the title “Master of Science”. To me science means the systematic analysis of the real or virtual world through observations and experiments as well as the further development of existing technology. I am lucky enough to be able to apply the knowledge and methodologies I learned during my studies to research projects at the Centre for Artificial Intelligence (CAI) of the Zurich University of Applied Sciences (ZHAW). I was mentored and supported during my studies by Prof. Dr. Thilo Stadelmann. He uses to say (also in accordance with his [blog-post](#)) “Great methodology delivers great theses”. It is always desirable to have an excellent outcome such as a system that can execute a task and thereby achieves or even overcomes state-of-the-art performance. However, in my opinion, it is equally or even more important to reason why and how something works, to justify choices, and to show limitations. I wrote my Thesis with these thoughts in mind and hope that the readers are able to follow my reasoning.

In the Introduction section, the fundamentals of Deep Learning and its limitations is described. Afterwards, it is motivated why methodologies inspired by neuroscience could overcome these limitations. This thesis aims at a target audience with a background in Deep Learning. Consequently, the concepts of Deep Learning are only roughly described. Since neurocomputing may be rather unknown to the target audience, a more extensive overview about this field is given.

I would like to thank a couple of colleagues and friends. First I think of my mentor Prof. Dr. Thilo Stadelmann who got me excited about AI years ago and later introduced me to research. He always encouraged creative ideas and helped to link different topics to address problems with methodologies from other fields. Thanks to his support, help, and guidance, I have grown personally as well as professionally. Further thanks go to Dr. Jan Deriu. Especially at the beginning of my thesis, he steered my thoughts in one direction. His unconventional thinking has led to the questioning of many methods that have stood the test of time for decades (this was also the inspiration for Geoffrey Hinton’s quote on the page before, although I wouldn’t presume to say that this thesis will change the future). To Prof. Dr. Christoph von der Malsburg for his seemingly endless patience in introducing me to neuroscience. He could build the bridge between the two diverging fields of Deep Learning and Neuroscience, serving as an inspiration for various new ideas.

The biggest thanks, however, goes to my family, who made this journey possible for me. My parents, who supported and encouraged me in every way. My younger brother who inspired me to study. My wife and son, who have been understanding and supportive and have always been the perfect counterbalance to the daily routine of studying. Without the support of my family, I would never have been able to embark on this academic path.



# Contents

<b>Zusammenfassung</b>	<b>vii</b>
<b>Abstract</b>	<b>ix</b>
<b>Preface</b>	<b>xi</b>
<b>Contents</b>	<b>xiii</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Fundamentals of Deep Learning . . . . .	1
1.2 Limitation of Deep Learning . . . . .	3
1.3 Differences to Biological Learning . . . . .	5
 <b>APPENDIX</b>	 <b>7</b>
<b>A TODO</b>	<b>9</b>
<b>Bibliography</b>	<b>11</b>



# List of Figures

1.1 Organization of the visual system in the cerebral cortex . . . . .	6
--	---

# List of Tables









Machine learning (ML) has become an indispensable part of our everyday lives. For example, we use it for machine translation, transport and logistics organization, product recommendations, fraud detection, self-driving cars, and much more. Machine Learning uses mathematical functions to map an input to an output. These functions usually extract patterns from the input data to build a relationship between input and output. The term Machine Learning stems from the fact that we use *machines* to correlate the input and the output to a function (i.e. to *learn* a function) during a training period. A sub-branch of Machine Learning is Deep Learning (DL). DL algorithms are able to learn hidden patterns within data to make predictions. They benefit from the accelerated computing power and big data made available in the last decade. Deep Learning is considered state-of-the-art for many learning tasks especially for high-dimensional data. Typical high-dimensional data are texts, audio recordings, 2D as well as 3D images, and videos. Deep Learning models use artificial neural networks to learn the mapping function between given input and output data. In the following, the fundamentals of Deep Learning is explained. Only those aspects that are relevant for the understanding of the rest of the thesis are discussed.

1.1 Fundamentals of Deep Learning	1
1.2 Limitation of Deep Learning	3
1.3 Differences to Biological Learning	5

## 1.1 Fundamentals

The idea for artificial neural networks (ANN) stems from biology and aims to capture the interaction of brain cells (neurons) with a mathematical model. A first model for a neuron was proposed by McCulloch and Pitts in 1943 [1]. Similar to how a neuron of the human brain transmits electrical impulses through the nervous system, the artificial neuron of McCulloch and Pitts receives multiple input signals and transforms them into a output signal. A neuron takes an input vector  $\mathbf{x} = (x_1, \dots, x_n)$  where  $x_i \in \{0, 1\}$  and maps it to an output  $\hat{y} \in \{0, 1\}$ . The mapping from the input to the output is done by using an aggregation function  $g$  that sums up the input vector  $\mathbf{x}$  and a activation function  $f$  that outputs 1 if the output of  $g$  is greater than a threshold  $\theta$  and 0 otherwise.

[1]: McCulloch et al. (1943)

$$g(\mathbf{x}) = g(x_1, \dots, x_n) = \sum_{i=1}^n x_i \quad (1.1)$$

$$\hat{y} = f(g(\mathbf{x})) = \begin{cases} 1, & \text{if } g(\mathbf{x}) \geq \theta \\ 0, & \text{otherwise} \end{cases} \quad (1.2)$$

In 1958, Rosenblatt [2] developed the Perceptron which works with real numbers as input. The input vector  $\mathbf{x} = (x_1, \dots, x_n)$  where  $x_i \in \mathbb{R}^n$  is multiplied with a weight vector  $\mathbf{w} = (w_1, \dots, w_n)$  where  $w_i \in \mathbb{R}^n$  with the same length.

[2]: Rosenblatt (1958)

$$g(\mathbf{x}) = g(x_1, \dots, x_n) = \sum_{i=1}^n w_i \cdot x_i \quad (1.3)$$

The output  $\hat{y} \in \{0, 1\}$  is similar to the McCulloch and Pitts neuron 1 if the aggregated value is greater than a threshold  $\theta$  and 0 otherwise as described in equation ?? . The equations (1.3) and (1.2) can be rewritten as

$$\hat{y} = f(g(\mathbf{x})) = \begin{cases} 1, & \text{if } \sum_{i=1}^n w_i \cdot x_i - \theta \geq 0 \\ 0, & \text{otherwise} \end{cases} \quad (1.4)$$

Later, the step-function  $f$  was replaced with other functions so that the output could also be a real number  $\hat{y} \in \mathbb{R}$ . Often used functions are

$$\begin{aligned} \text{Sigmoid: } \sigma(z) &= \frac{1}{1 + e^{-z}} \\ \text{Rectified linear unit (ReLU): } (z)^+ &= \max(0, z) \\ \text{Hyperbolic tangent (tanh): } \tanh(z) &= \frac{e^z - e^{-z}}{e^z + e^{-z}} \end{aligned} \quad (1.5)$$

By convention, instead of using the threshold  $-\theta$  often a bias  $b$  is used which leads to:

$$\begin{aligned} z = g(\mathbf{x}) &= \mathbf{w} \cdot \mathbf{x} = \sum_{i=1}^n w_i \cdot x_i + b \\ \hat{y} &= f(z) \end{aligned} \quad (1.6)$$

However, the brain consists of multiple neurons which are connected through synapses. Therefore, ANNs consist not only of one neuron but combine multiple neurons in a network. These neurons are organized into layers. A shallow neural network consists of one hidden layer in which the input  $\mathbf{x}$  is fed to calculate the output  $\hat{y}$ . The universal approximation theorem of Cybenko [3] proves that a shallow network with enough neurons can approximate any mapping function between inputs and outputs. However, very complex mapping functions may need too many hidden neurons. The neurons of the hidden layer extract features in the input space. Since only one layer is used, the features cannot be hierarchically organized and become complex enough.

[3]: Cybenko (1989)

A multi-layer perceptron (MLP), on the other hand, consists of multiple layers. The different layers extract increasingly complex features. In a MLP, the input  $\mathbf{x}$  is fed into the first layer, each subsequent layer  $l$  gets the output of the previous layer  $l - 1$  as input. In the following, we describe the mathematical model for fully connected layers, where all neurons of a layer are connected with the subsequent layer. For a MLP with  $m$  layers, we define the output of the aggregation  $g$  as  $\mathbf{z}^{[l]}$  and the output the activation function as  $\mathbf{a}^{[l]}$  for layer  $l$ . Furthermore, we use  $\mathbf{w}^{[l]}$  for the weight vector and  $b^{[l]}$  for the bias of layer  $l$ . Thus, the mathematical model of a MLP is defined as

$$\begin{aligned} z^{[l]} &= w^{[l]} a^{[l-1]} + b^{[l]} \\ a^{[l]} &= f(z^{[l]}) \end{aligned} \quad (1.7)$$

Since the input is fed into the first layer and the output is the result from the last layer  $x = a^{[0]}$  and  $\hat{y} = a^{[m]}$  holds true.

So far, only the forward-pass which is used to calculate the output  $\hat{y}$  was discussed. However, the model output  $\hat{y}$  will only be close to the target output  $y$  if the weights  $w^{[l]}$  and biases  $b^{[l]}$  are properly defined. These parameters are learned during a training period. The training can take place in a supervised, semi-supervised, self-supervised, unsupervised, or reinforcement learning based manner. In supervised learning, the output of the model  $\hat{y}$  for a given input  $x$  is compared to manually created target outputs  $y$ . Unsupervised learning, on the other hand, tries to find patterns in the input  $x$  and to cluster the samples into meaningful groups without using target labels. Semi-supervised learning is a hybrid approach of the the aforementioned principles that combines a small amount of labelled data with a large amount of unlabelled data. In self-supervised learning, the target outputs  $y$  are directly derived from the input data  $x$  (e.g. predict a masked part of the input  $x$ ). Lastly, reinforcement learning algorithms aim to maximize a reward that they become from an environment based on some action they executed.

These learning principle have in common that a loss function  $\mathcal{L}$  can calculate a loss value based on the model output  $\hat{y}$ . For example, the mean square error (MSE) can be used for regression problems or the negative log-likelihood for classification problems. The chosen loss function is minimized iteratively with stochastic gradient descent (SGD). The idea behind stochastic gradient descent is to make use of the fact that the negative gradient of the loss value points to the direction of the steepest descent (i.e. in the direction where the loss gets smaller). SGD therefore updates the network parameters by taking a step of size  $\eta$  in the direction of their negative gradient

$$\begin{aligned} \Delta w^{[l]} &= -\eta \nabla_{w^{[l]}} \mathcal{L} \\ \Delta b^{[l]} &= -\eta \nabla_{b^{[l]}} \mathcal{L} \end{aligned} \quad (1.8)$$

The gradients of the weights  $w^{[l]}$  and biases  $b^{[l]}$  can efficiently be calculated with an algorithm called backpropagation [4], which is just a smart implementation of the chain rule

[4]: Rumelhart et al. (1986)

TODO: Describe CNN, Transformer, ... etc.???

## 1.2 Limitations of Deep Learning

The rise of Deep Learning over the past decade has only been possible because of major technological advances in hardware. Moore's law [5] states that the number of transistors in a dense integrated circuit doubles about every two years and is the only known physical process following an exponential curve. An analysis by OpenAI shows that since 2012 the amount of compute has even increasing exponentially with a doubling

time of 3.4 months [6]. However, the exponential increase seems to come to an end since the size of transistors hit physical limitations. It is assumed that Moore's law will end by around 2025 [7]. Besides the progress in the field, Deep Learning methods also became better because they grew exponentially. They not only have more layers and more parameters but also require more data. Even the growth in the last five years is astonishing. While the state-of-the-art language model from 2018 [8] had around 94M parameters, the state-of-the-art in 2020 [9] already had 175B parameters. Training such a model on a single V100 GPU would take about 355 years and cost about 4.6M dollars [10]. A recent language model from Microsoft and Nvidia [11] even has 530B parameters. Only a few institutions with massive resources are able to train such big models. In general, inference on low-budget hardware such as smartphones or embedded hardware becomes prohibitive with the growing size of deep networks. Even though there exist techniques to shrink the model size after training such as quantization [12], model pruning [13], or model distillation [14] it is questionable if making models bigger is the best way to develop intelligent systems.

Another major issue of Deep Learning systems is that they suffer from catastrophic forgetting. If a model is trained on a specific task and afterwards trained (or fine-tuned) on another task, the model suffers a "catastrophic" drop in performance over the first task. The reason for this effect is that the model during training on the second task adjusts the parameters learned during the first task and therefore "forgets" the learned mapping functions. Just mixing all datasets or to learn all tasks in parallel in a current multi-task setup [15] doesn't seem feasible to achieve some kind of general intelligence as this involves too many different unrelated tasks. Catastrophic forgetting is also caused by the fact that learning is mostly done offline. Online learning [16] and lifelong learning [17] are currently hot research topics. However, these methods have not yet been established.

Furthermore, there exist problems which may cannot be solved with the current principles used for Deep Learning. First of all, it is questionable if Deep Learning models can achieve *real* generalization. With enough data, can achieve generalization in the sense that the model can interpolate within the data distribution. However, deep learning models fail to extrapolate. For example, convolutional neural networks (CNNs) do not generalize to different viewpoints unless they are added to the training data [18].

Second, Deep Learning is not able to learn abstract relationships in a few trials but requires many samples of it and is thus data hungry.<sup>1</sup> Marcus Gary [19] argues that if he tells that a "schmister" is a sister over the age of 10 but under the age of 21, humans can immediately infer whether they or their best friends have any "schmister". However, modern DL systems lack a mechanism for learning abstractions through explicit, verbal definition and require thousands or even more training samples.

Third, no DL model has been able to demonstrate causal reasoning in a generic way. Deep Learning models find correlations between the inputs and the outputs, but not the causation. Other AI approaches such as hierarchical Bayesian computing or probabilistic graphical models

[7]: Kumar (2015)

[8]: Peters et al. (2018)

[9]: Brown et al. (2020)

[11]: Shoenberger et al. (2020)

[18]: Madan et al. (2022)

[19]: Marcus (2018)

are better at causal reasoning but cannot be well combined with Deep Learning models.

Lastly, Deep Learning models are to some extent too isolated since they have no embodiment and cannot interact with the world. For example, the human body provides needs, goals, emotions, and gut feeling. In current Deep Learning systems, emotions are totally absent and the goals are set externally. Deep Reinforcement Learning can be considered as a first step in the direction of dissolving this isolation, as they interact with a virtual environment. AI systems that interact with the real world do not work well so far. Moravec's paradox [20] states that "it is comparatively easy to make computers exhibit adult level performance on intelligence tests or playing checkers, and difficult or impossible to give them the skills of a one-year-old when it comes to perception and mobility".

[20]: Moravec (1995)

## 1.3 Biological Learning

The human brain comprises many interconnected areas processing everything in parallel. For example, Figure 1.1 illustrates the connections between different organizational units in the cerebral cortex which are responsible for vision. It can be seen that these areas are connected in a rather complex structure. Deep Learning architectures, on the other hand, are mostly unidirectional and the signal flows unidirectional from layer to layer. However, the choice of the architecture influences how the model can learn the mapping function from input to output. It could be that the complex structure of our brain comprises an inductive bias which was learned over time through evolution.

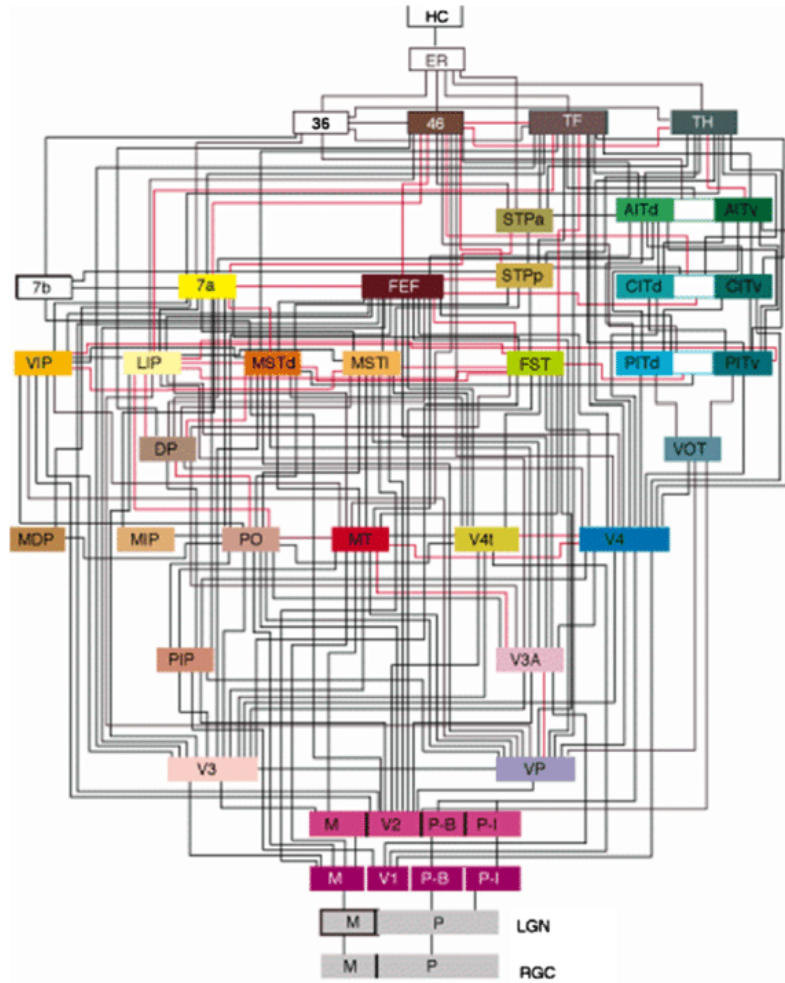
A learning system requires a mechanism that tells the system if something goes well or wrong so that it can learn from it. This is called the *credit assignment problem*. Backpropagation (c.f. Section 1.1) solves this problem by propagating the error backwards through the network. However, information flows in the brain only in one direction from the presynaptic neurons to the postsynaptic neurons. Therefore, backpropagation is not biologically plausible. Lillicrap et al. [22] shows that an additional set of random feedback weights is able to transmit useful gradients. Their work has reopened questions how the brain could process error signals and has dispelled some long-held assumptions about algorithmic constraints on learning.

[22]: Lillicrap et al. (2016)

Not only the structure of the network and the way how the feedback is calculated is different between biological learning and Deep Learning. Also the neurons themselves are different. While the artificial neuron doesn't have any dynamics (c.f. Equation (1.6)), biological neurons are highly dynamic: Biological neurons adapt their firing rate to constant inputs, they may continue firing after an input disappears, and can even fire when no input is active.

TODO: Add reference to reservoir computing

Lastly, the neurons in the brain are self-organizing. This means that a group of elementary units such as neurons or a group of neurons perform similar rule of behavior on a sub-set of the available information. Such a system doesn't have a central supervision that orchestrates these



**Figure 1.1:** The organization of the visual system in the cerebral cortex. The image is from Felleman et al. [21].

units. Each unit applies similar deterministic functions to the information received. Two important principles of such systems are (i) localized learning which means that each unit adapt their behavior to the information they receive; and (ii) emergence which means that there is no explicit loss function that tells the system what to do.



# APPENDIX



TODO | **A**

TODO



# Bibliography

Here is the list of references in citation order.

- [1] Warren S. McCulloch and Walter Pitts. 'A logical calculus of the ideas immanent in nervous activity'. en. In: *The Bulletin of Mathematical Biophysics* 5.4 (Dec. 1943), pp. 115–133. doi: [10.1007/BF02478259](#) (cited on page 1).
- [2] F. Rosenblatt. 'The perceptron: A probabilistic model for information storage and organization in the brain.' en. In: *Psychological Review* 65.6 (1958), pp. 386–408. doi: [10.1037/h0042519](#) (cited on page 1).
- [3] G. Cybenko. 'Approximation by superpositions of a sigmoidal function'. en. In: *Mathematics of Control, Signals, and Systems* 2.4 (Dec. 1989), pp. 303–314. doi: [10.1007/BF02551274](#) (cited on page 2).
- [4] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. 'Learning representations by back-propagating errors'. en. In: *Nature* 323.6088 (Oct. 1986), pp. 533–536. doi: [10.1038/323533a0](#) (cited on page 3).
- [5] Gordon E. Moore. 'Cramming more components onto integrated circuits, Reprinted from Electronics, volume 38, number 8, April 19, 1965, pp.114 ff.' In: *IEEE Solid-State Circuits Society Newsletter* 11.3 (Sept. 2006), pp. 33–35. doi: [10.1109/N-SSC.2006.4785860](#) (cited on page 3).
- [6] Open AI. *AI and Compute*. 2018. URL: <https://openai.com/blog/ai-and-compute/> (visited on 08/19/2022) (cited on page 4).
- [7] Suhas Kumar. 'Fundamental Limits to Moore's Law'. In: arXiv:1511.05956 (Nov. 2015). arXiv:1511.05956 [cond-mat] (cited on page 4).
- [8] Matthew Peters et al. 'Deep Contextualized Word Representations'. en. In: *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*. New Orleans, Louisiana: Association for Computational Linguistics, 2018, pp. 2227–2237. doi: [10.18653/v1/N18-1202](#) (cited on page 4).
- [9] Tom Brown et al. 'Language Models are Few-Shot Learners'. In: *Advances in Neural Information Processing Systems*. Ed. by H. Larochelle et al. Vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901 (cited on page 4).
- [10] Lambda. *OpenAI's GPT-3 Language Model: A Technical Overview*. 2021. URL: <https://lambdalabs.com/blog/demystifying-gpt-3/> (visited on 08/19/2022) (cited on page 4).
- [11] Mohammad Shoeybi et al. 'Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism'. In: arXiv:1909.08053 (Mar. 2020). arXiv:1909.08053 [cs] (cited on page 4).
- [12] Hao Wu et al. 'Integer Quantization for Deep Learning Inference: Principles and Empirical Evaluation'. In: arXiv:2004.09602 (Apr. 2020). arXiv:2004.09602 [cs, stat] (cited on page 4).
- [13] Tejalal Choudhary et al. 'A comprehensive survey on model compression and acceleration'. en. In: *Artificial Intelligence Review* 53.7 (Oct. 2020), pp. 5113–5155. doi: [10.1007/s10462-020-09816-7](#) (cited on page 4).
- [14] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 'Distilling the Knowledge in a Neural Network'. In: arXiv:1503.02531 (Mar. 2015). arXiv:1503.02531 [cs, stat] (cited on page 4).
- [15] Yu Zhang and Qiang Yang. 'A Survey on Multi-Task Learning'. In: arXiv:1707.08114 (Mar. 2021). arXiv:1707.08114 [cs] (cited on page 4).
- [16] Doyen Sahoo et al. 'Online Deep Learning: Learning Deep Neural Networks on the Fly'. In: arXiv:1711.03705 (Nov. 2017). arXiv:1711.03705 [cs] (cited on page 4).
- [17] German I. Parisi et al. 'Continual lifelong learning with neural networks: A review'. en. In: *Neural Networks* 113 (May 2019), pp. 54–71. doi: [10.1016/j.neunet.2019.01.012](#) (cited on page 4).

- [18] Spandan Madan et al. 'When and how convolutional neural networks generalize to out-of-distribution category–viewpoint combinations'. en. In: *Nature Machine Intelligence* 4.2 (Feb. 2022), pp. 146–153. doi: [10.1038/s42256-021-00437-5](https://doi.org/10.1038/s42256-021-00437-5) (cited on page 4).
- [19] Gary Marcus. 'Deep Learning: A Critical Appraisal'. In: arXiv:1801.00631 (Jan. 2018). arXiv:1801.00631 [cs, stat] (cited on page 4).
- [20] Hans Moravec. *Mind children: the future of robot and human intelligence*. eng. 4. print. Cambridge: Harvard Univ. Press, 1995 (cited on page 5).
- [21] D. J. Felleman and D. C. Van Essen. 'Distributed Hierarchical Processing in the Primate Cerebral Cortex'. en. In: *Cerebral Cortex* 1.1 (Jan. 1991), pp. 1–47. doi: [10.1093/cercor/1.1.1](https://doi.org/10.1093/cercor/1.1.1) (cited on page 6).
- [22] Timothy P. Lillicrap et al. 'Random synaptic feedback weights support error backpropagation for deep learning'. en. In: *Nature Communications* 7.1 (Dec. 2016), p. 13276. doi: [10.1038/ncomms13276](https://doi.org/10.1038/ncomms13276) (cited on page 5).

